Training began at Dec 16, 2023, 2:24:46 PM and is still in progress.

Status	Training
Training pipeline ID	7579113076813725696
Created	Dec 16, 2023, 2:24:46 PM
Start time	Dec 16, 2023, 2:24:46 PM
Budget (original)	1 node hours
Elapsed time	1 hr 15 min
Region	us-central1
Encryption type	Google-managed

Dataset	magawa
Target column	target
Data split	Randomly assigned (80/10/10)
Column metadata	VIEW DETAILS

Algorithm	AutoML
Objective	Tabular classification
Optimized for	AUC ROC
Training stage	Training

Training performance

Training performance

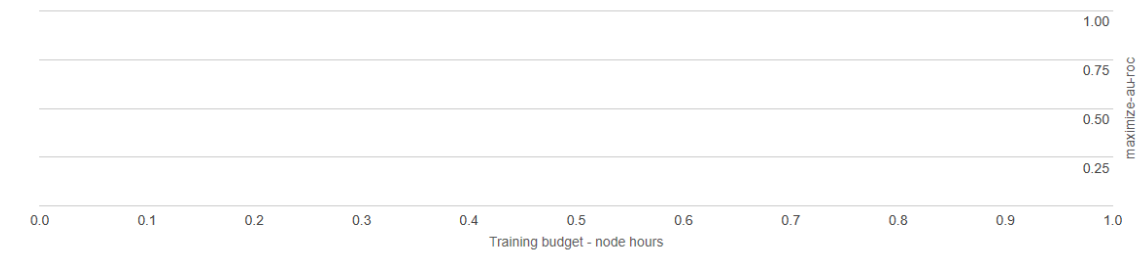

Training

TRAIN NEW MODEL | REFRESH | LEARN

TRAINING PIPELINES | CUSTOM JOBS | HYPERPARAMETER TUNING JOBS | NAS JOBS

Training pipelines are the primary model training workflow in Vertex AI. You can use training pipelines to create an AutoML-trained model or a custom-trained model. For custom-trained models, training pipelines orchestrate custom training jobs and hyperparameter tuning with additional steps like adding a dataset or uploading the model to Vertex AI for prediction serving. Learn more

Region: us-central1 (Iowa)

Filter: Enter a property name

Name	ID	Status	Job type	Model type	Duration	Last updated	Created	Ended	Labels
magawa	7579113076813725696	Finished	Training pipeline	Tabular classification	2 hr 53 min	Dec 16, 2023, 5:18:19 PM	Dec 16, 2023, 2:24:46 PM	Dec 16, 2023, 5:18:19 PM	–

magawa > Version 1 | VIEW DATASET | EXPORT

EVALUATE | DEPLOY & TEST | BATCH PREDICT | VERSION DETAILS

untitled_2830597194215490743 | OVERVIEW | SLICE ANALYSIS | COMPARE | CREATE EVALUATION

Labels

Filter

All labels	0.996
0	0.997
1	0.077

Evaluation details

Confidence threshold: 0.5

All labels

PR AUC	0.996
ROC AUC	0.996
Log loss	0.058
F1 score	0.9879593
Micro-F1	0.9879593
Macro-F1	0.4969716
Precision	98.8%
Recall	98.8%

To evaluate your model, set the **confidence threshold** to see how precision and recall are affected. The best confidence threshold depends on your use case. Read some example scenarios to learn how evaluation metrics can be used.

Precision-recall curve | ROC curve | Precision-recall by threshold

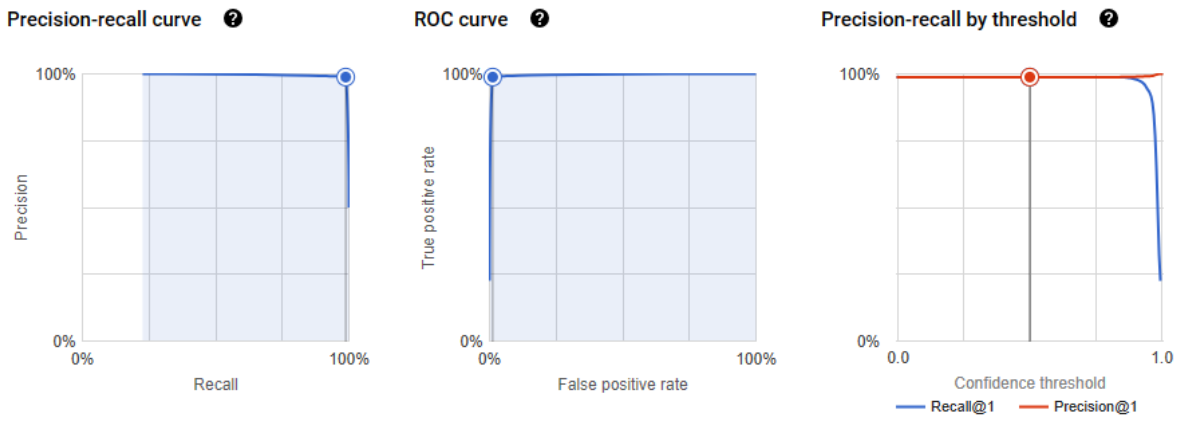

Confusion matrix

This table shows how often the model classified each label correctly (in blue), and which labels were most often confused for that label (in gray).

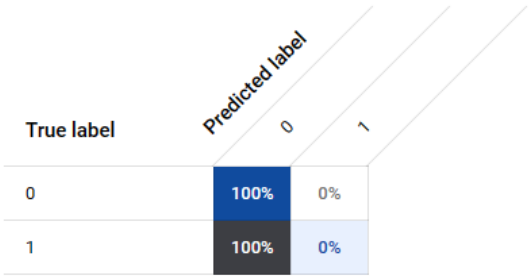

Feature importance

Model feature attribution tells you how important each feature is when making a prediction. Attribution values are expressed as a percentage; the higher the percentage, the more strongly that feature impacts a prediction on average. Model feature attribution is expressed using the Sampled Shapley method. Learn more

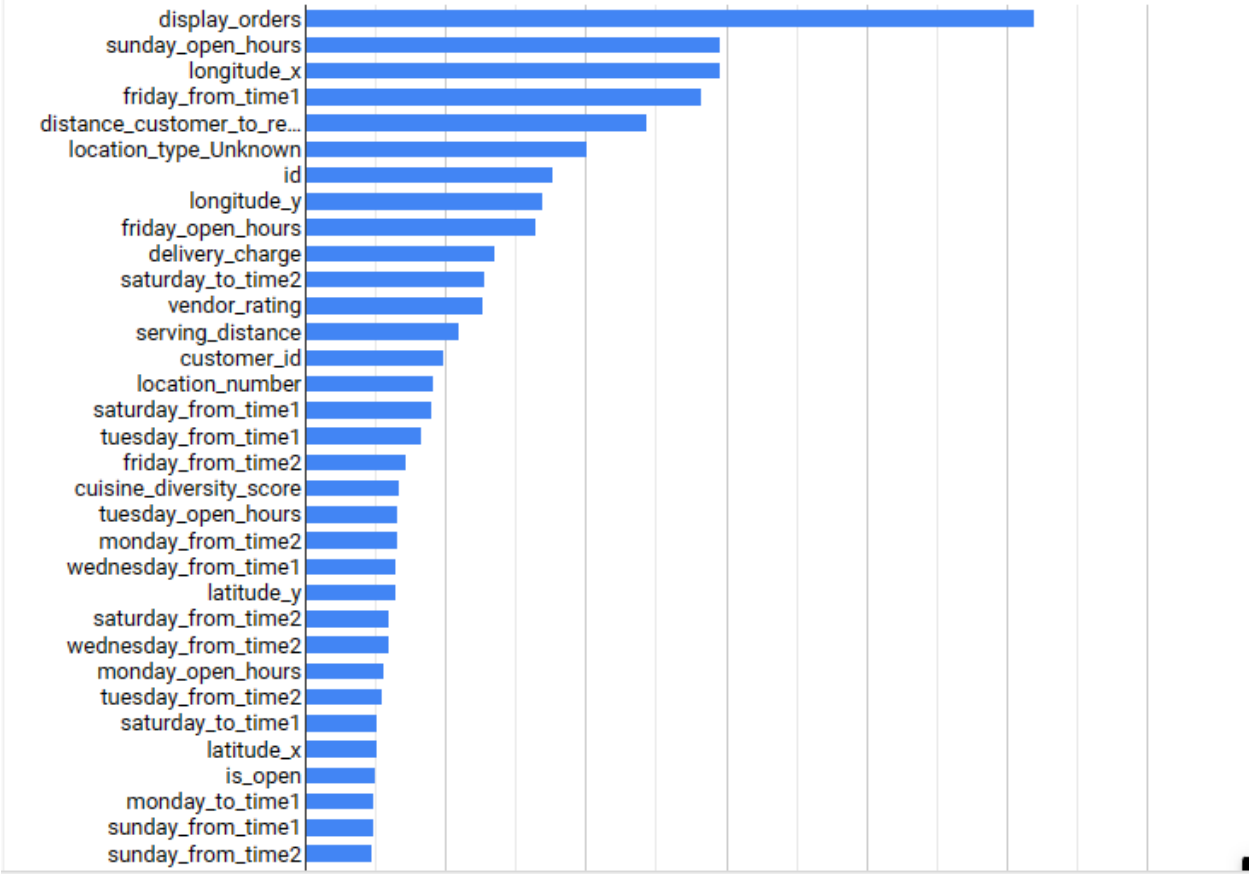

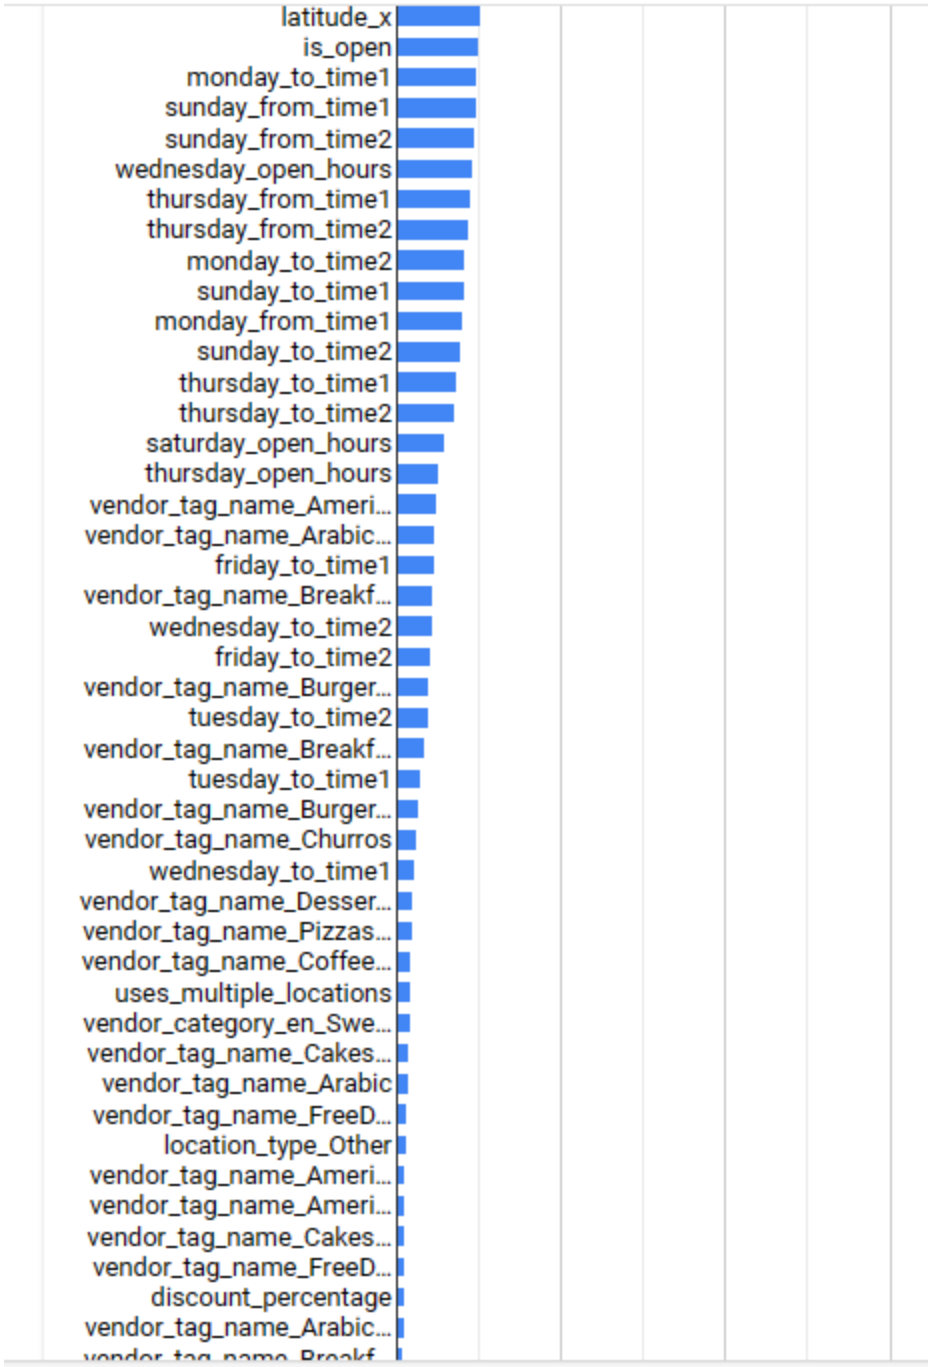
latitude_x
is_open
monday_to_time1
sunday_from_time1
sunday_from_time2
wednesday_open_hours
thursday_from_time1
thursday_from_time2
monday_to_time2
sunday_to_time1
monday_from_time1
sunday_to_time2
thursday_to_time1
thursday_to_time2
saturday_open_hours
thursday_open_hours
vendor_tag_name_Ameri...
vendor_tag_name_Arabic...
friday_to_time1
vendor_tag_name_Breakf...
wednesday_to_time2
friday_to_time2
vendor_tag_name_Burger...
tuesday_to_time2
vendor_tag_name_Breakf...
tuesday_to_time1
vendor_tag_name_Burger...
vendor_tag_name_Churros
wednesday_to_time1
vendor_tag_name_Desser...
vendor_tag_name_Pizzas...
vendor_tag_name_Coffee...
uses_multiple_locations
vendor_category_en_Swe...
vendor_tag_name_Cakes...
vendor_tag_name_Arabic
vendor_tag_name_FreeD...
location_type_Other
vendor_tag_name_Ameri...
vendor_tag_name_Ameri...
vendor_tag_name_Cakes...
vendor_tag_name_FreeD...
discount_percentage
vendor_tag_name_Arabic...

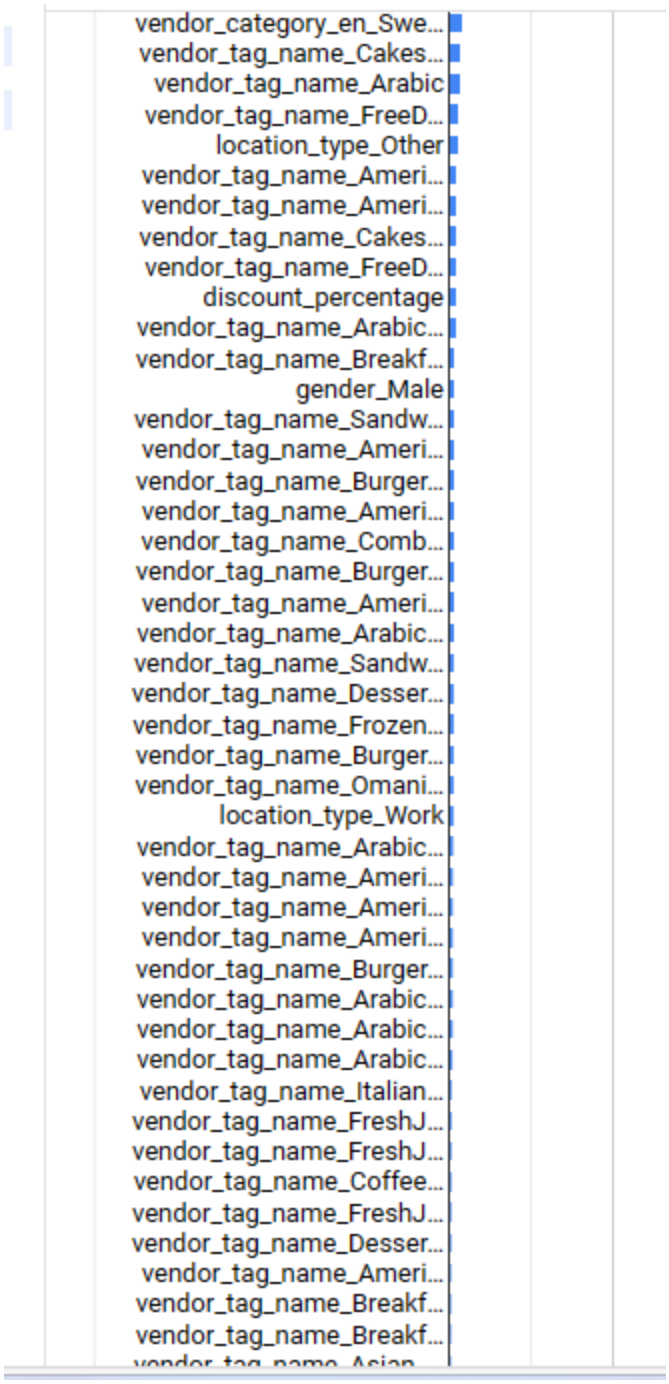
vendor_category_en_Swe...
vendor_tag_name_Cakes...
vendor_tag_name_Arabic
vendor_tag_name_FreeD...
location_type_Other
vendor_tag_name_Ameri...
vendor_tag_name_Ameri...
vendor_tag_name_Cakes...
vendor_tag_name_FreeD...
discount_percentage
vendor_tag_name_Arabic...
vendor_tag_name_Breakf...
gender_Male
vendor_tag_name_Sandw...
vendor_tag_name_Ameri...
vendor_tag_name_Burger...
vendor_tag_name_Ameri...
vendor_tag_name_Comb...
vendor_tag_name_Burger...
vendor_tag_name_Ameri...
vendor_tag_name_Arabic...
vendor_tag_name_Sandw...
vendor_tag_name_Desser...
vendor_tag_name_Frozen...
vendor_tag_name_Burger...
vendor_tag_name_Omani...
location_type_Work
vendor_tag_name_Arabic...
vendor_tag_name_Ameri...
vendor_tag_name_Ameri...
vendor_tag_name_Ameri...
vendor_tag_name_Burger...
vendor_tag_name_Arabic...
vendor_tag_name_Arabic...
vendor_tag_name_Arabic...
vendor_tag_name_Italian...
vendor_tag_name_FreshJ...
vendor_tag_name_FreshJ...
vendor_tag_name_Coffee...
vendor_tag_name_FreshJ...
vendor_tag_name_Desser...
vendor_tag_name_Ameri...
vendor_tag_name_Breakf...
vendor_tag_name_Breakf...

Compare evaluations PREVIEW

You can compare up to 5 evaluations of the same type (regression, classification, etc.).

magawa > Version 1 > untitled_2830597194215490743 ADD

Confidence threshold 0.5

Model version	Evaluation	PR AUC	ROC AUC	Log loss	F1 score	F1-Micro	F1-Macro	Precision	Recall
Version 1	untitled_2830597194215490743	0.996	0.996	0.058	0.988	0.988	0.497	98.8%	98.8%

Precision-recall curve

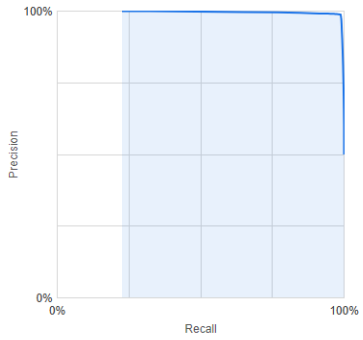

ROC curve

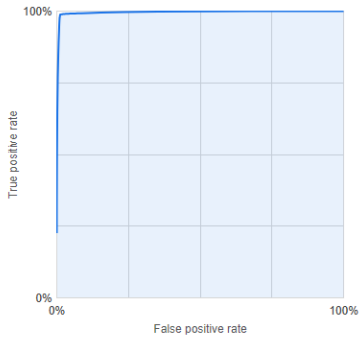

Precision-recall by threshold

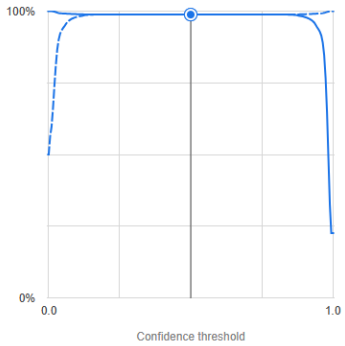

Feature importance

Compare evaluations PREVIEW

Feature importance

Model feature attribution tells you how important each feature is when making a prediction. Attribution values are expressed as a percentage; the higher the percentage, the more strongly that feature impacts a prediction on average. Model feature attribution is expressed using the Sampled Shapley method. Learn more

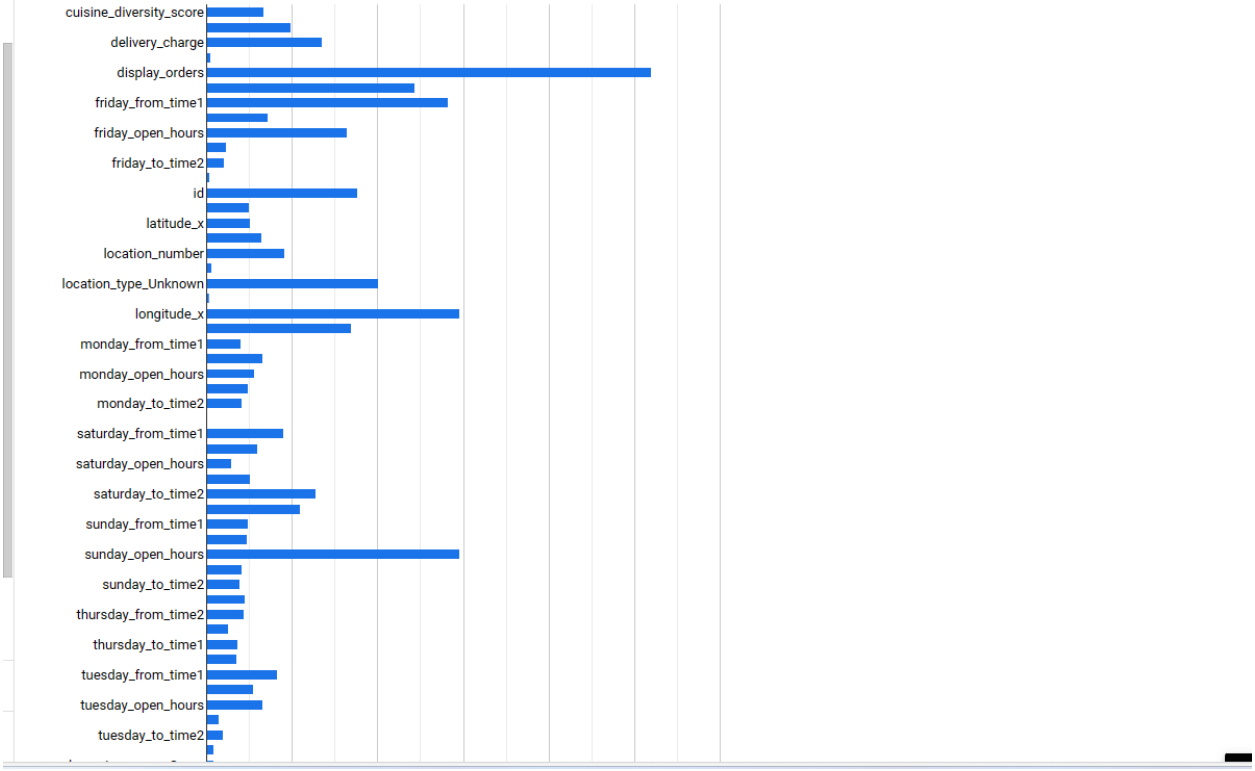

magawa > Version 1 VIEW DATASET EXPORT

EVALUATE DEPLOY & TEST BATCH PREDICT VERSION DETAILS

Use your edge-optimized model

Container

Export your model as a TF Saved Model to run on a Docker container.

Deploy your model

Endpoints are machine learning models made available for online prediction requests. Endpoints are useful for timely predictions from many users (for example, in response to an application request). You can also request batch predictions if you don't need immediate results.

DEPLOY TO ENDPOINT

Name	ID	Status	Models	Deployment resource pool	Region	Monitoring	Most recent monitoring job	Most recent alerts	Last updated	API	Labels	Encryption
No active endpoints containing this model												

Test your model PREVIEW

In order to test your model, you will need to deploy it first. Pricing guide

Deploy to endpoint

1. Define your endpoint
2. Model settings
3. Model monitoring
4. Monitoring objectives

DEPLOY CANCEL

magawa (Version 1)

Traffic split *
100 %

Compute resources

Choose how compute resources will serve prediction traffic to your model

- **Autoscaling**: If you set a minimum and maximum, compute nodes will scale to meet traffic demand within those boundaries
- **No scaling**: If you only set a minimum, then that number of compute nodes will always run regardless of traffic demand (the maximum will be set to minimum)

Once scaling settings are set, they can't be changed unless you redeploy the model. Pricing guide

Minimum number of compute nodes *
1

Default is 1. If set to 1 or more, then compute resources will continuously run even without traffic demand. This can increase cost but avoid dropped requests due to node initialization.

Maximum number of compute nodes (optional)
1

Enter a number equal to or greater than the minimum nodes. Can reduce costs but may cause reliability issues for high traffic.

ADVANCED SCALING OPTIONS

Machine type *
n1-standard-8, 8 vCPUs, 30 GiB memory

Logging

Logging settings are permanent for this endpoint, and Logging charges will apply. To change your logging preference in the future, create a new endpoint. Learn more

- [x] Enable access logging for this endpoint
- [ ] Disable container logging for this endpoint

Explainability options

- (•) No explainability
- ( ) Feature attribution
- ( ) Example-based explanation

ADVANCED SCALING OPTIONS

Machine type *
n1-standard-8, 8 vCPUs, 30 GiB memory

Logging

Logging settings are permanent for this endpoint, and Logging charges will apply. To change your logging preference in the future, create a new endpoint. Learn more

- [x] Enable access logging for this endpoint
- [ ] Disable container logging for this endpoint

Explainability options

- (•) No explainability
- ( ) Feature attribution
- ( ) Example-based explanation

It may take several minutes for endpoint settings to take effect.

DONE

ADD A MODEL

CONTINUE

Deploy your model

Endpoints are machine learning models made available for online prediction requests. Endpoints are useful for timely predictions from many users (for example, in response to an application request). You can also request batch predictions if you don't need immediate results.

DEPLOY TO ENDPOINT

Name	ID	Status	Models	Deployment resource pool	Region	Monitoring	Most recent monitoring job	Most recent alerts	Last updated	API	Labels
Restaurant Recommender	7552406832728244224	Deploying model	0	–	us-central1	–	–	–	Dec 16, 2023, 7:12:53 PM	SAMPLE REQUEST	

Test your model PREVIEW

Your model must be successfully deployed to an endpoint before you can test it.

Feature column name	Type	Value	Local feature importance
customer_id	Text	YYKZN0F	–
location_number	Text	0.0	–
latitude_x	Text	0.6957402667917951	–
longitude_x	Text	0.706456720971424	–

Predicted column not yet known

Prediction result

–

Deploy your model

Endpoints are machine learning models made available for online prediction requests. Endpoints are useful for timely predictions from many users (for example, in response to an application request). You can also request batch predictions if you don't need immediate results.

DEPLOY TO ENDPOINT

Name	ID	Status	Models	Deployment resource pool	Region	Monitoring	Most recent monitoring job	Most recent alerts	Last updated	API	Labels	Encryption
Restaurant Recommender	7552406832728244224	Active	1	–	us-central1	Disabled	–	–	Dec 16, 2023, 7:30:04 PM	SAMPLE REQUEST		Google-managed

Test your model PREVIEW